

**SPRINGONE2GX**

WASHINGTON, DC

# Call Tracing

Monish Unni

[monish.unni@gmail.com](mailto:monish.unni@gmail.com)



springone *ZBX*



# About Me

Monish Unni  
Programmer | Architect  
Java | Spring | Scala | NodeJS

# Contents

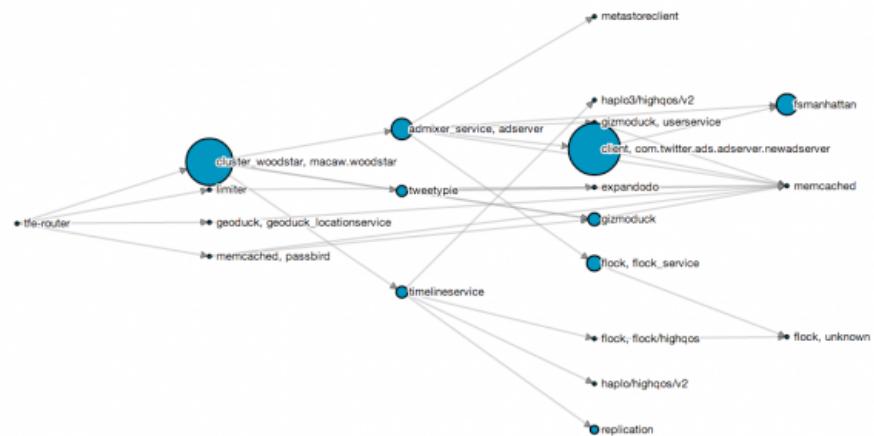
- Background
- Use case
- Core Concepts
- Demo
- Minimal Tracing
- Zipkin Tracing
- Demo
- Production Deployment
- Implementing suggestions
- Exceptions to the model
- Lessons Learnt

# Background

- In the beginning we had
  - Clearly specified contracts and roles
  - Well known topology and flow
    - web servers, application servers, services, databases
  - Prescheduled down times
  - Stacktraces were good enough
    - for the most part.

# Background

- Today we have microservices that are a
  - Collection of software modules
  - Developed by different teams
  - Different programming languages
  - Span many machines
  - Multiple physical facilities



# Background

- Developers have checked in and tested code
- SRE monitors services, queues and databases
  - queue surges
  - database tracing
  - disk corruption
  - load re-balancing

# Background

- SRE needs to
  - shutdown and redeploy services
  - map a service to other services
  - map a service to database/tables
  - figure out weak links in a system
  - figure out critical paths

# Background

- SRE needs to
  - debug the system
  - figure out where the bottlenecks are?
  - need to determine if request paths are correct?
  - correlate call tree with system behavior.

- In April 2010, Google published a research paper called as “Dapper”

Modern Internet services are often implemented as complex, large-scale distributed systems. These applications are constructed from collections of software modules that may be developed by different teams, perhaps in different programming languages, and could span many thousands of machines across multiple physical facilities.

-- Google Dapper Paper

- Using an in-house dapper-like framework, we are able to layer the request demand through the aggregate infrastructure onto a simple visualization
  - Netflix Feb 2015<sup>1</sup>
- ..help us gather timing data for all the disparate services involved in managing a request to the Twitter API...
  - Twitter June 2012<sup>2</sup>
- Real-time distributed tracing and call graphs have proven to be indispensable for insight into the service performance of web applications, and the associated cost across the entire technology stack at LinkedIn.
  - LinkedIn Feb 2015<sup>3</sup>
- Talk on “Call Tracing”
  - SpringOne2GX Sept 2015

<sup>1</sup> [http://techblog.netflix.com/2015\\_02\\_01\\_archive.html](http://techblog.netflix.com/2015_02_01_archive.html)

<sup>2</sup> <https://blog.twitter.com/2012/distributed-systems-tracing-with-zipkin>

<sup>3</sup> <https://engineering.linkedin.com/distributed-service-call-graph/real-time-distributed-tracing-website-performance-and-efficiency>



Jason Johnson

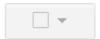
5

+ Share



Services ?

Gmail



More

1-20 of 20



COMPOSE

Inbox

All Mail

Trash

Circles

Family

Friends

Fun



Search people...

Jason

Jeremy

Michael

Miles

Peter

Bali photos: <https://...>

Brittany

Elizabeth

Nathania

Paul

<input type="checkbox"/> Vickie	Events Brainstorm for Fall 2012-13 - Cape Cod HHKS board retreat Dim sum t...	Jun 19
<input type="checkbox"/> Anissa Mak	Parents' Day - Anissa and Daniel want to say Happy Parents' Day to r...	Jun 19
<input type="checkbox"/> Leslie	Trip to Vegas - Which...	
<input type="checkbox"/> Puiyan, me (3)	Apartment hunting - h...	
<input type="checkbox"/> Peter, me (2)	Home-cooked Dinner...	
<input type="checkbox"/> Miles, me (2)	Hangout? - ...	
<input type="checkbox"/> Linda, me (2)	Victoria's bask...	
<input type="checkbox"/> Kat	Froyo + catch-up time...	
<input type="checkbox"/> Jordan	House...	
<input type="checkbox"/> Meian	Grandma's b...	
<input type="checkbox"/> Zinnia	Hey :) - I miss you! W...	
<input type="checkbox"/> me, Brett (2)	Team c...	
<input type="checkbox"/> Shal, me (2)	Go for a run next Tue...	
<input type="checkbox"/> Meredith Blackwell	congratulations!! - He...	
<input type="checkbox"/> Leslie, me (3)	BBQ by the Ch...	
<input type="checkbox"/> Anissa Mak	Did you see the colou...	

This weekend

To   Cc Bcc

This weekend

Hey guys,

Thanks for the help with the trip last weekend! Here's a shot of the view from the cabin!



Send



+



Google

Gmail

Compose

Inbox

All Mail

Trash

Circles

Family

Friends

Fun

Friends

Fun

Holiday

Victoria's basket

Kat

Jordan

Meian

Zinnia

me, Brett (2)

Shal, me (2)

Meredith Blackwell

Leslie, me (3)

Anissa Mak

Vickie

Anissa Mak

Leslie

Puiyan, me (3)

Peter, me (2)

Miles, me (2)

Linda, me (2)

Kat

Jordan

Meian

Zinnia

me, Brett (2)

Shal, me (2)

Meredith Blackwell

Leslie, me (3)

Anissa Mak

Events Brainstorm for Fall 2012-13 - Cape Cod HHKS board retreat Dim sum to Jun 19

Family Parents' Day Anissa and Daniel want to say Happy Parents' Day to r Jun 19

Trip to Vegas - Which Apartment hunting - H

Housewarming Dinner

Hey guys,

Thanks for the help with the trip last weekend! Here's a shot of the view from the cabin!

Send

Did you see the solar

This weekend

To Steven Bills x Phil Sharp x Cc Bcc

Jun 19

1–20 of 20 < > ⚙

Unless otherwise indicated, these slides are © 2013-2015 Pivotal Software, Inc. and licensed under a Creative Commons Attribution-NonCommercial license: <http://creativecommons.org/licenses/by-nc/3.0/>

Services (guess) .

Contacts –

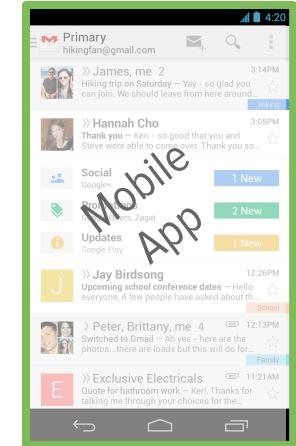
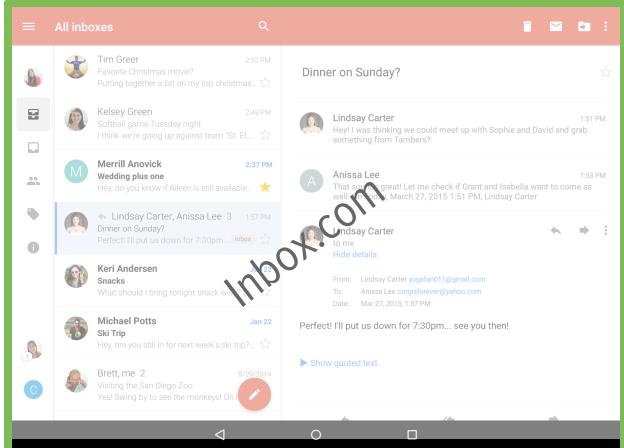
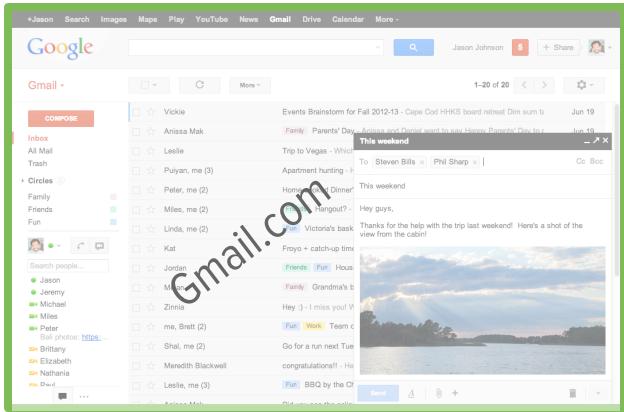
Presence –

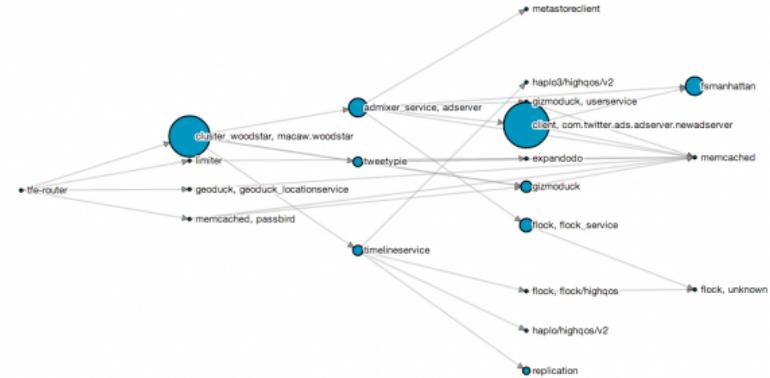
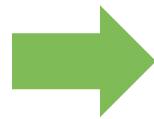
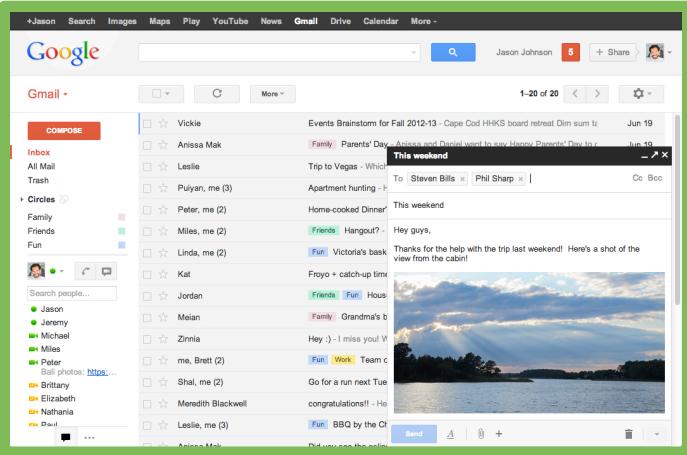
Circles –

Hangout audio –

Hangout video –

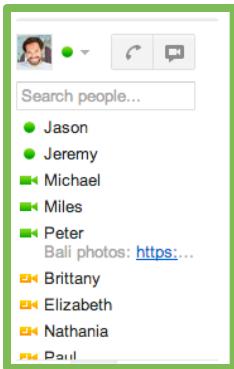
Labels/Tags –

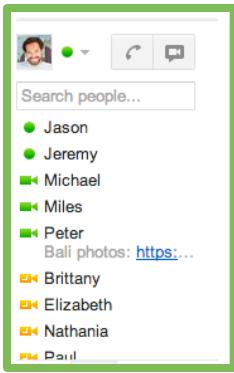


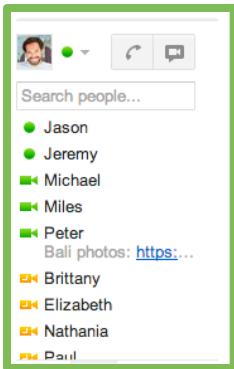


A single request can run across 100s/1000s of machines/services and involve hundreds of different subsystems.

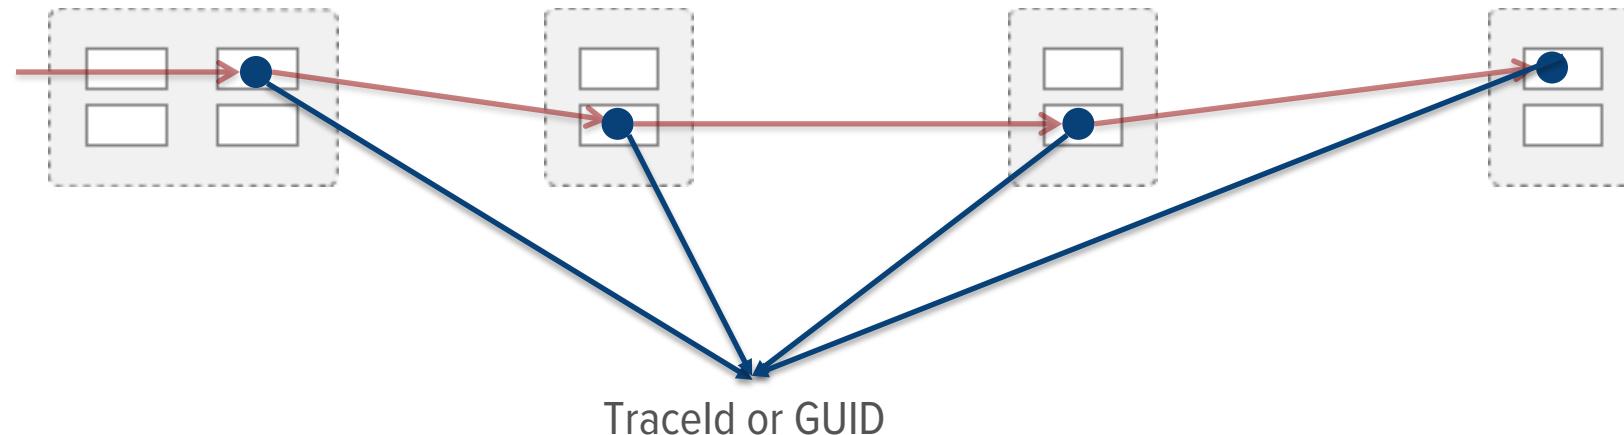
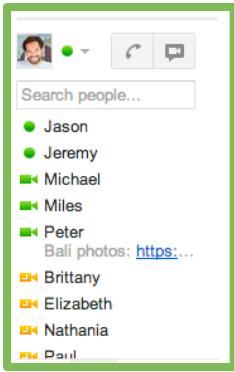
# Call Tracing Concepts

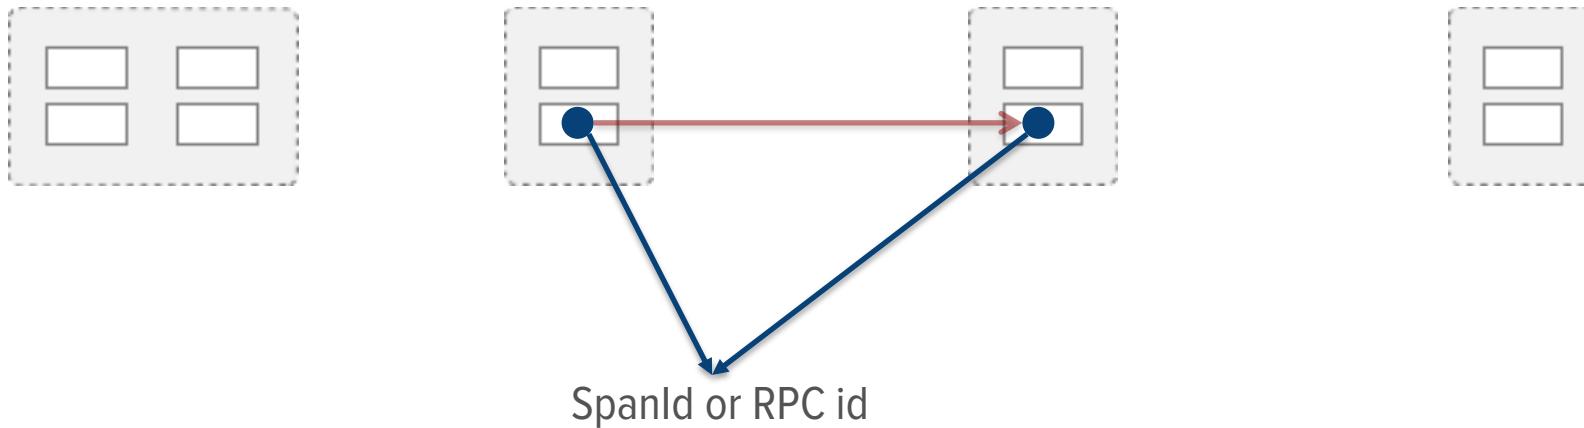
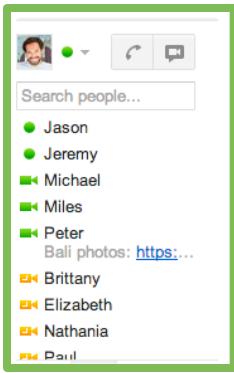






Traceld or GUID



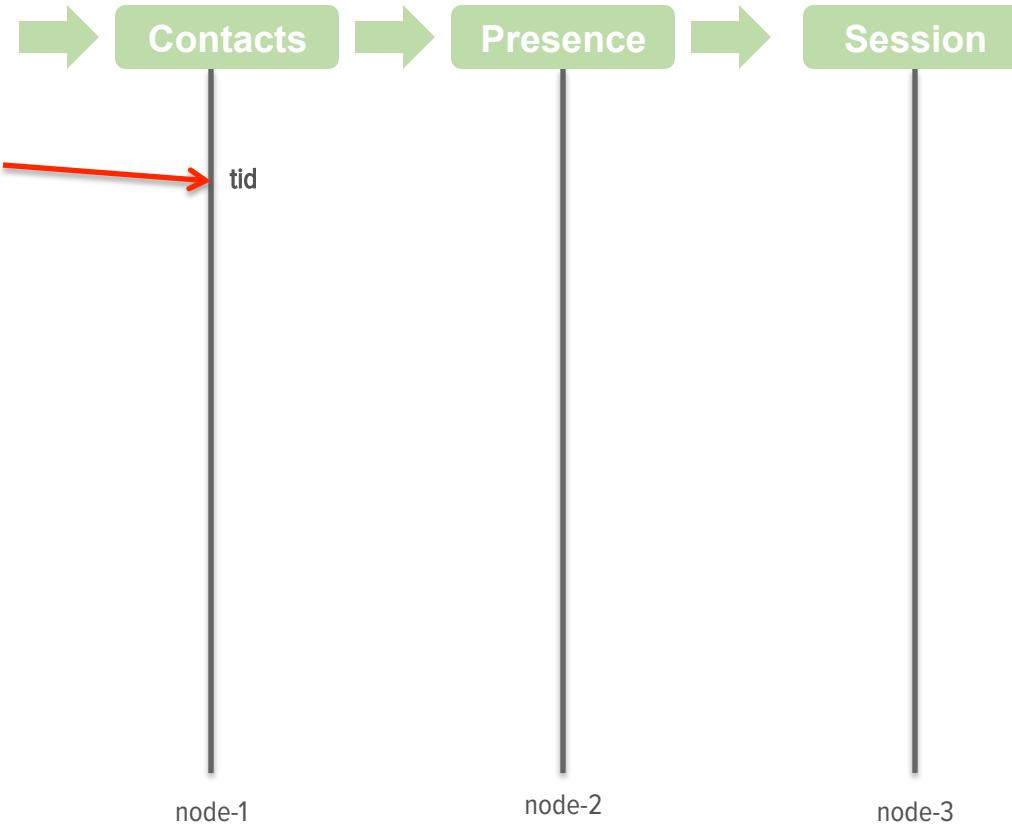
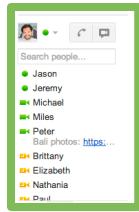


# Review

- Traceld
  - a globally unique immutable Id
  - correlates to a request
- SpanId
  - a unique id for an RPC

# Demo

- Start two services
  - server\_a, server\_b
- client calls server\_a
  - server\_a calls server\_b
- show trace UX



## Goal

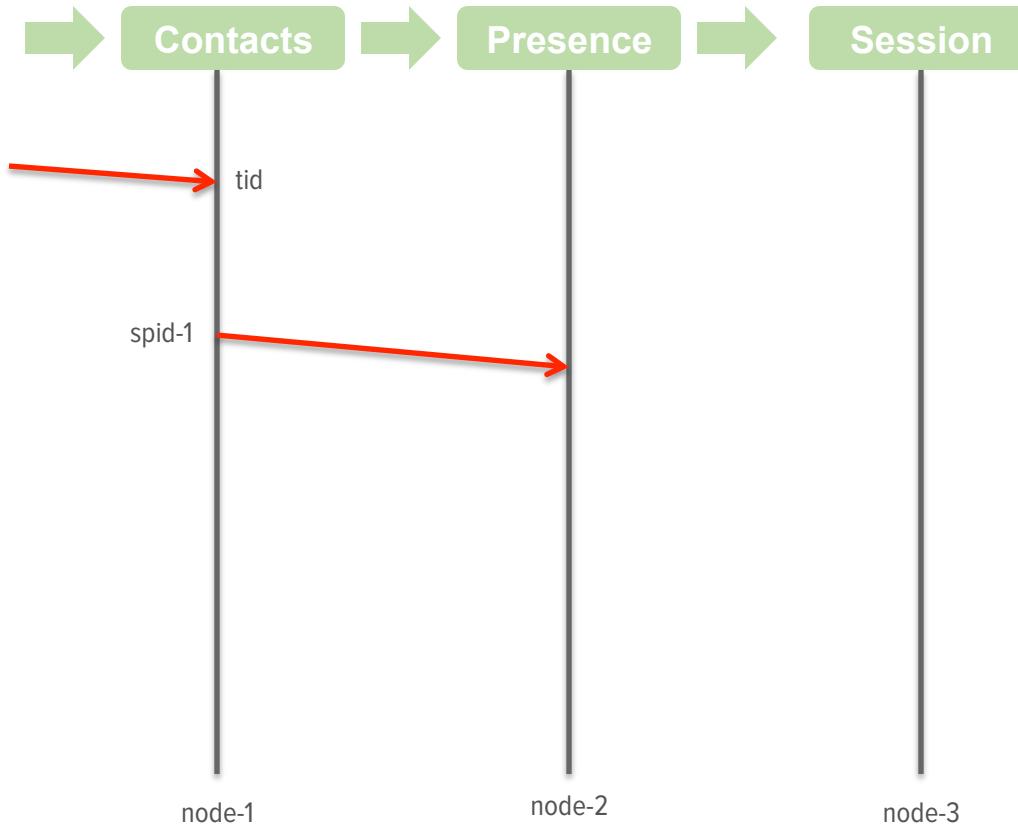
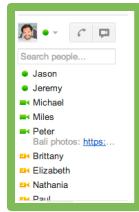
Stitch call tree from local logs

## Node Log View

**node-1:** tid

**node-2:**

**node-3:**



## Goal

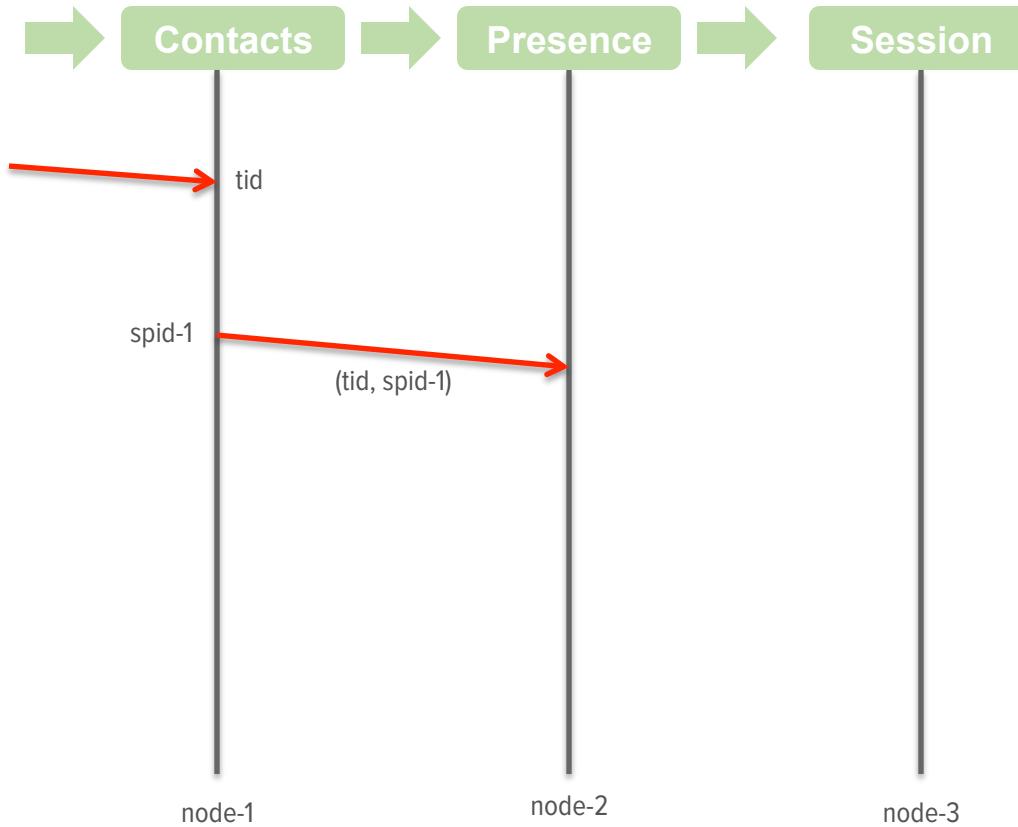
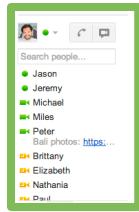
Stitch call tree from local logs

## Node Log View

**node-1:** tid, spid-1

**node-2:**

**node-3:**



## Goal

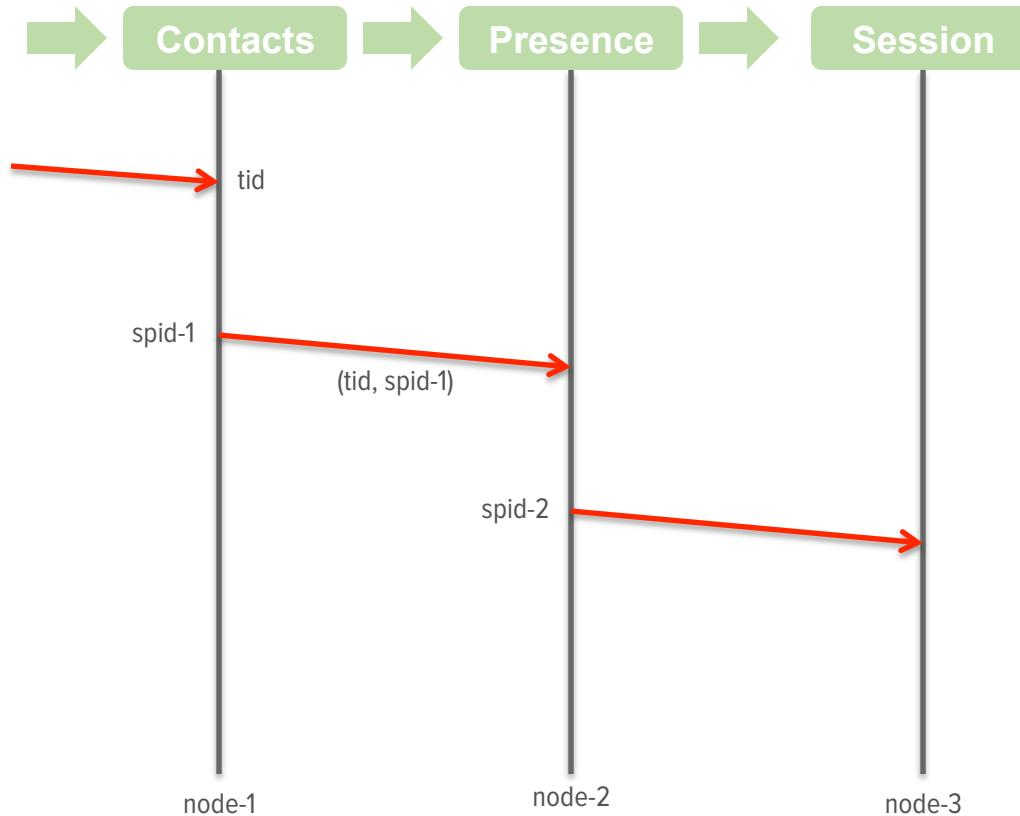
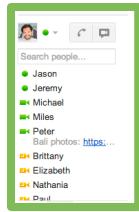
Stitch call tree from local logs

## Node Log View

**node-1:** tid, spid-1

**node-2:** tid, spid-1

**node-3:**



## Goal

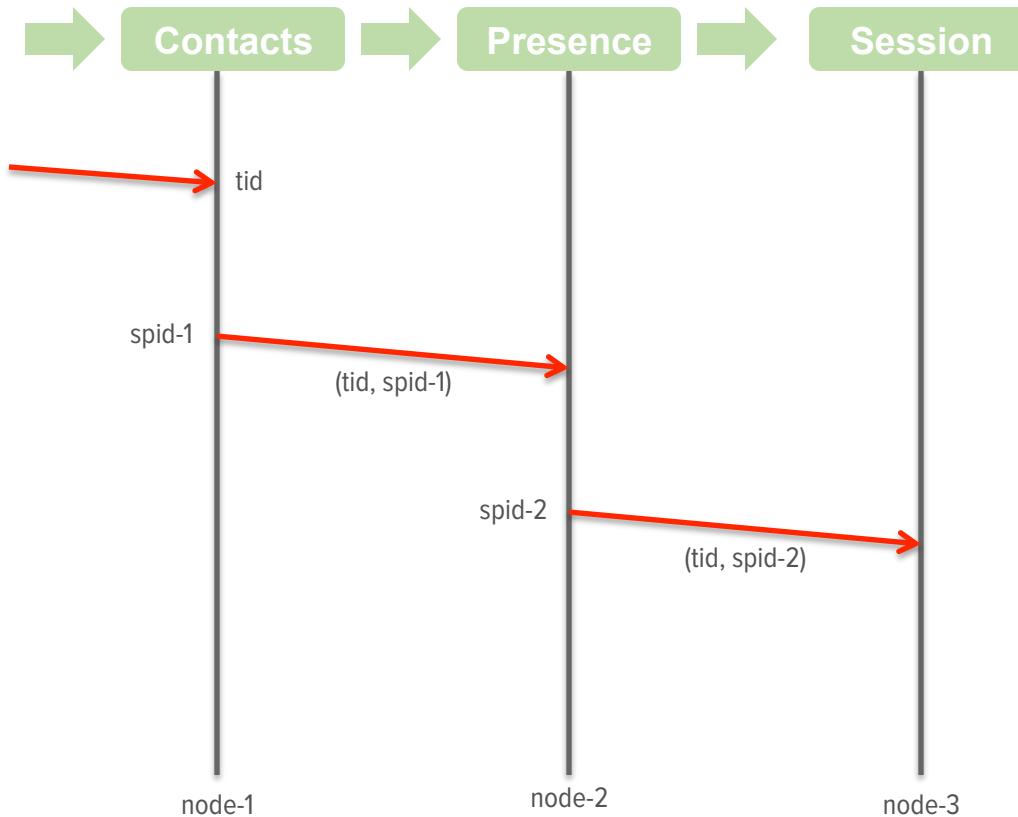
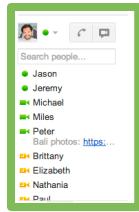
Stitch call tree from local logs

## Node Log View

**node-1:** tid, spid-1

**node-2:** tid, spid-1, spid-2

**node-3:**



## Goal

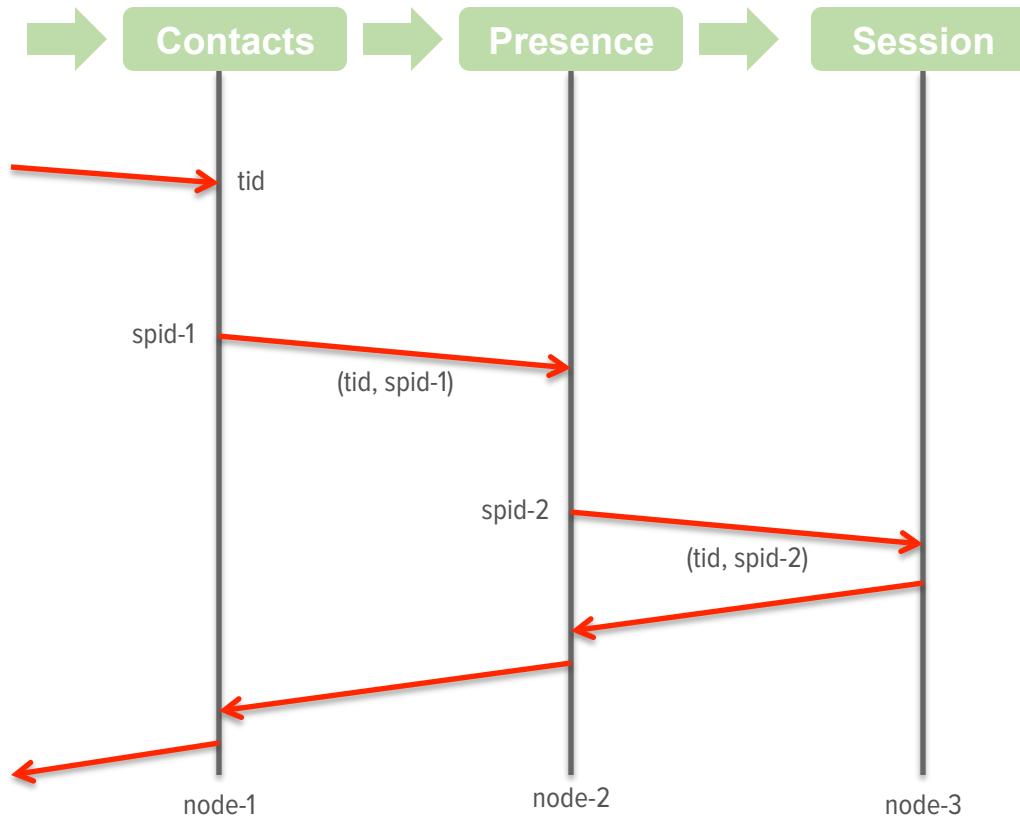
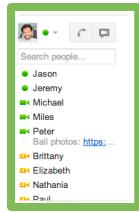
Stitch call tree from local logs

## Node Log View

**node-1:** tid, spid-1

**node-2:** tid, spid-1, spid-2

**node-3:** tid, spid-2



## Goal

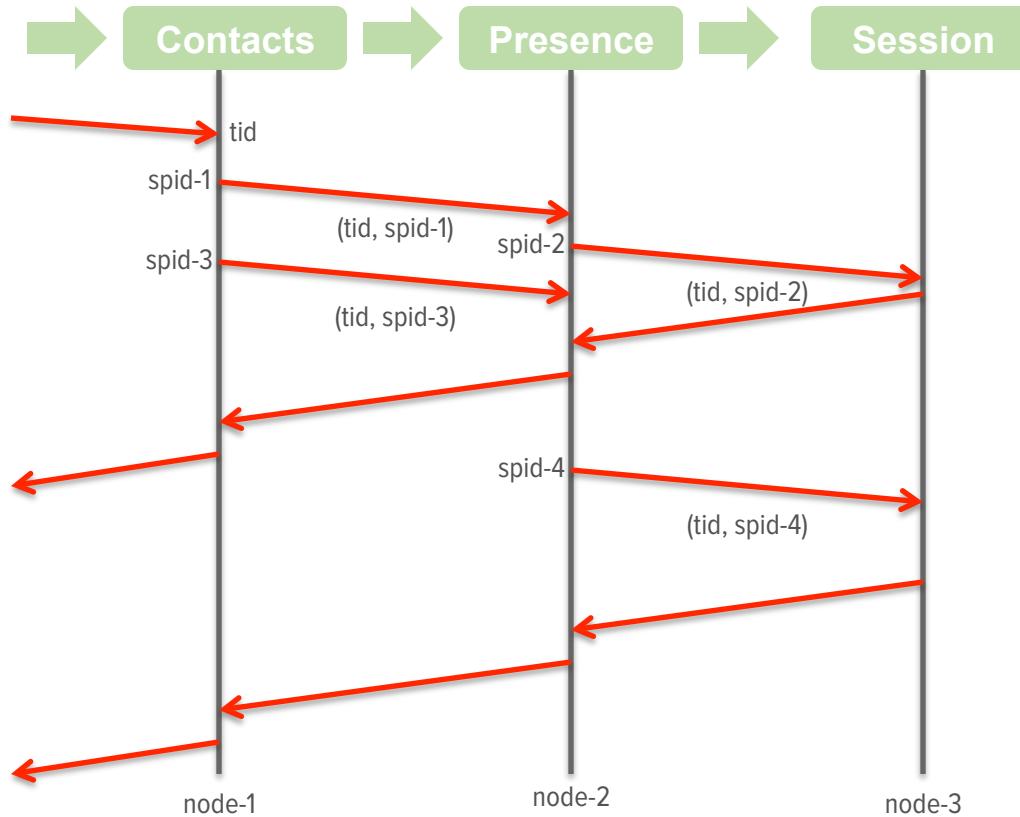
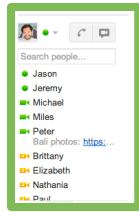
Stitch call tree from local logs

## Node Log View

**node-1:** tid, spid-1

**node-2:** tid, spid-1, spid-2

**node-3:** tid, spid-2



## Goal

Stitch call tree from local logs

## Node Log View

**node-1**: tid, spid-1, spid-3

**node-2**: tid, spid-1, spid-2, spid-3

**node-3**: tid, spid-2, spid-4

## Limit

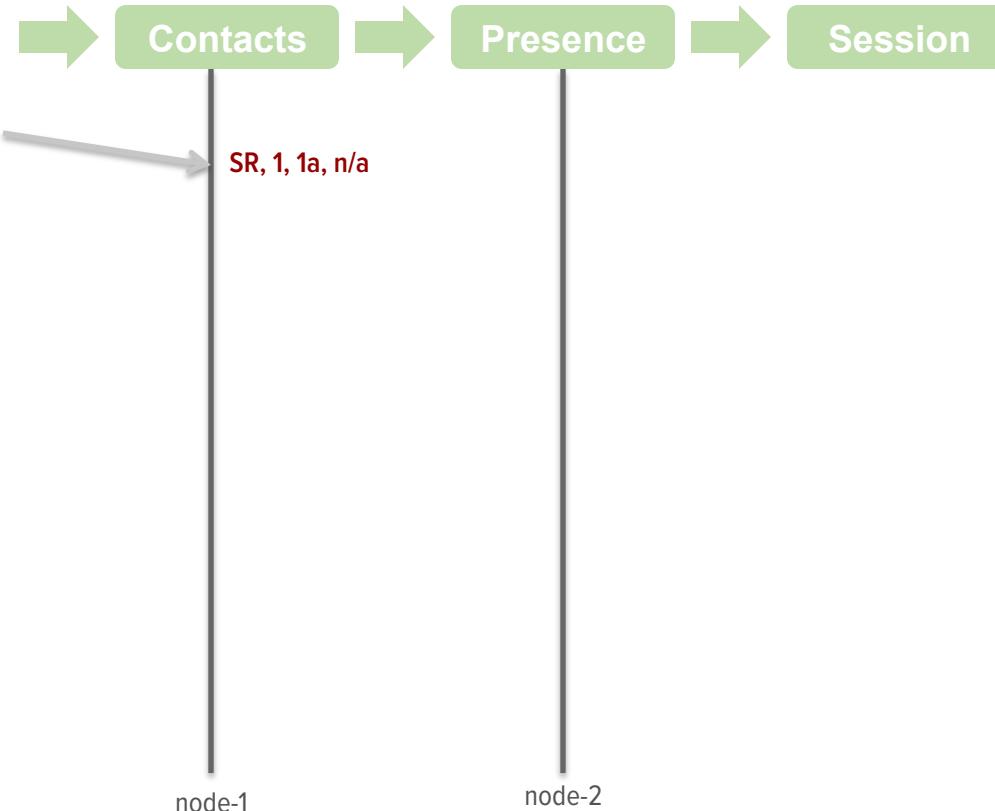
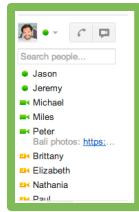
- can correlate all requests
- cannot stitch call tree from logs

# Review

- Black-box instrumentation
  - Pros
    - Application level transparency
    - Traceld, SpanId, Timestamp allow correlation of requests
  - Cons
    - Request Paths or call trees require
      - reasoning based on statistical inference
      - derivation of causal relationships

# Uniform Logging Structure

- Traceld, SpanId, ParentId, Event +Timestamp
- Four key events with timestamps
  - Server Receive/Send (**SR/SS**)
    - duration for a service call
  - Client Send/Receive (**CS/CR**)
    - duration for a client call
- ParentId
  - Optional – root span has none
  - Previous SpanId



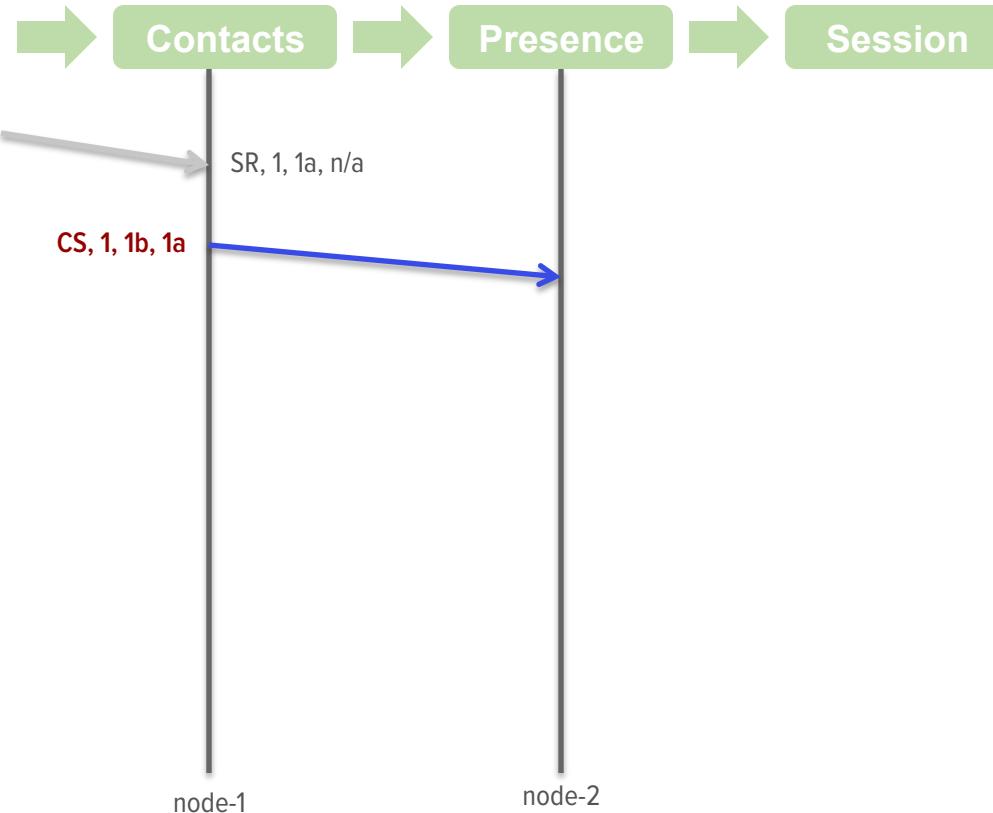
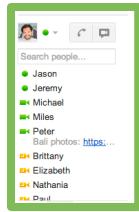
## Structure

Event: **Server Received (SR)**

TracId: 1

SpanId: 1a

ParentId : N/A



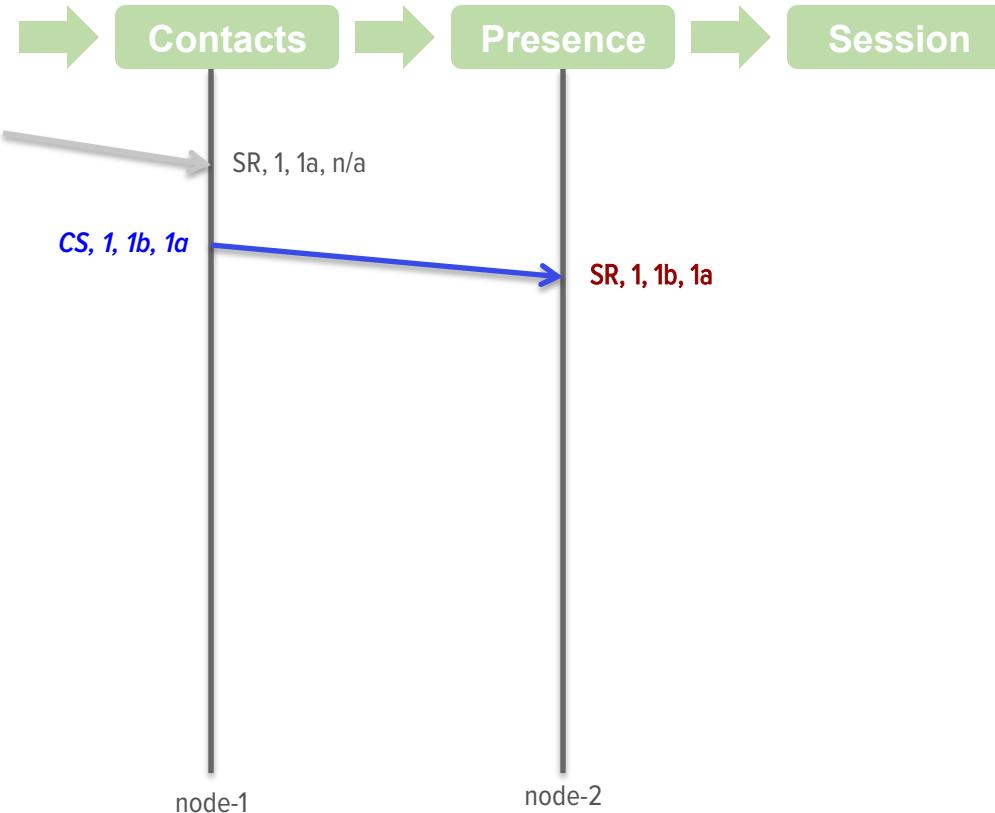
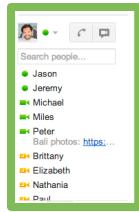
## Structure

Event: **Client Sent (CS)**

TracId: 1

SpanId: 1b

ParentId : 1a



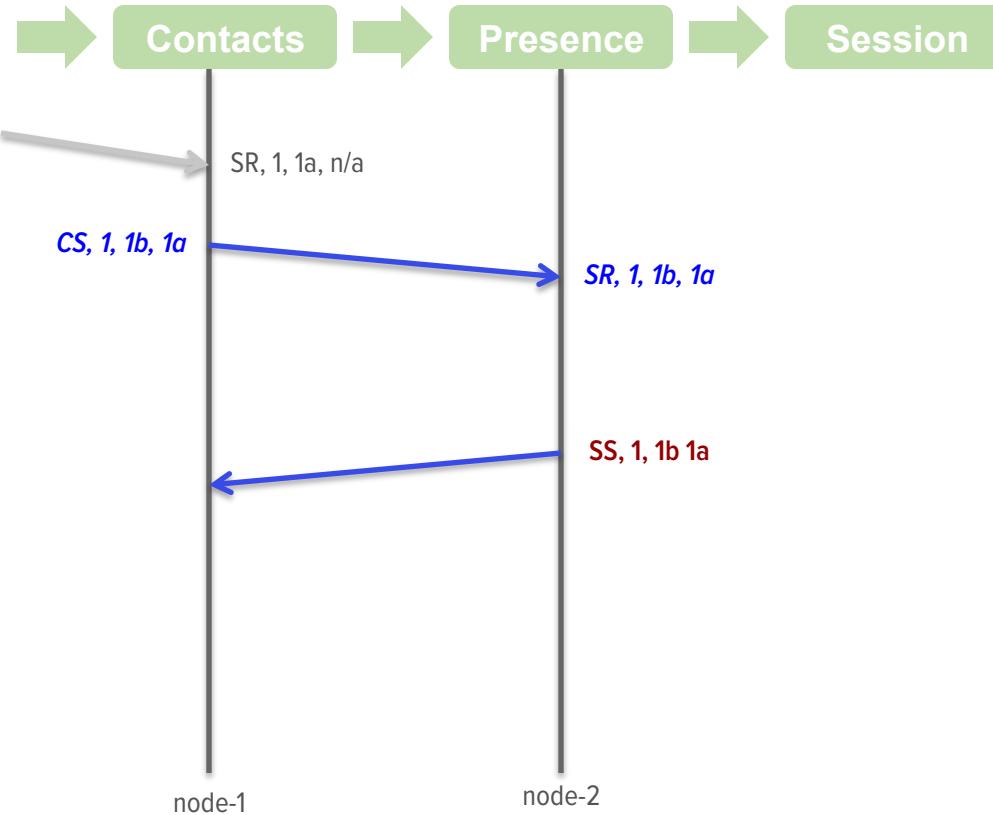
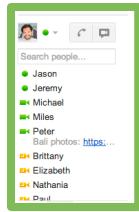
## Structure

Event: **Server Received(SR)**

TracId: 1

SpanId: 1b

ParentId : 1a



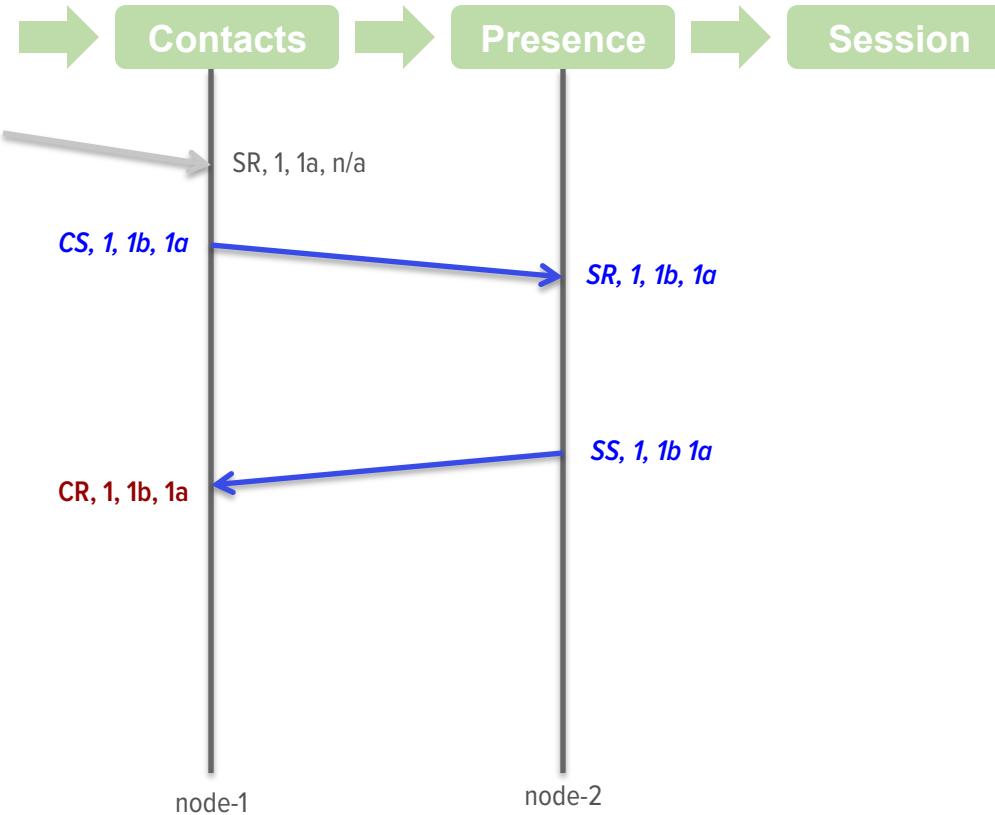
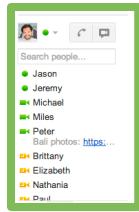
## Structure

Event: **Server Sent(SS)**

Traceld: 1

SpanId: 1b

ParentId : 1a



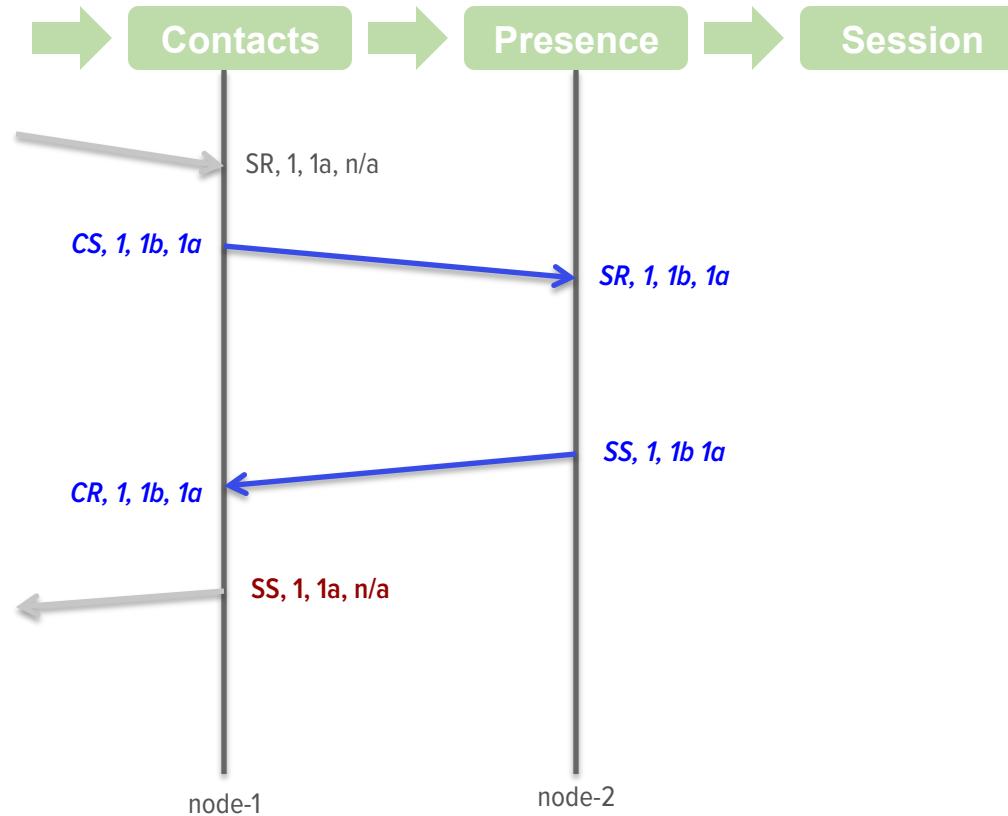
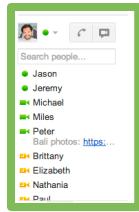
## Structure

Event: Client Received(CR)

TracId: 1

SpanId: 1b

ParentId : 1a



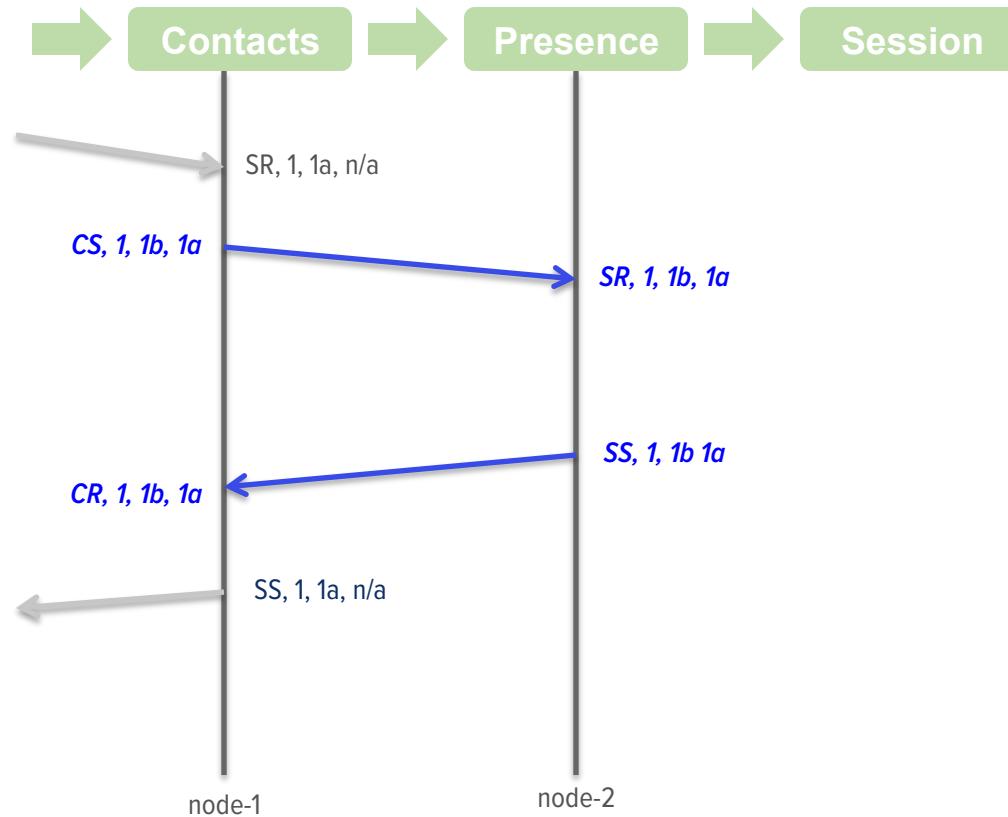
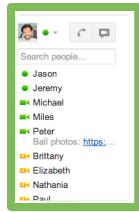
## Structure

Event: **Server Sent(SS)**

Traceld: 1

SpanId: 1a

ParentId : N/A



Span

CS, SR, SS, CR

Traceid : 1

SpanId : 1b

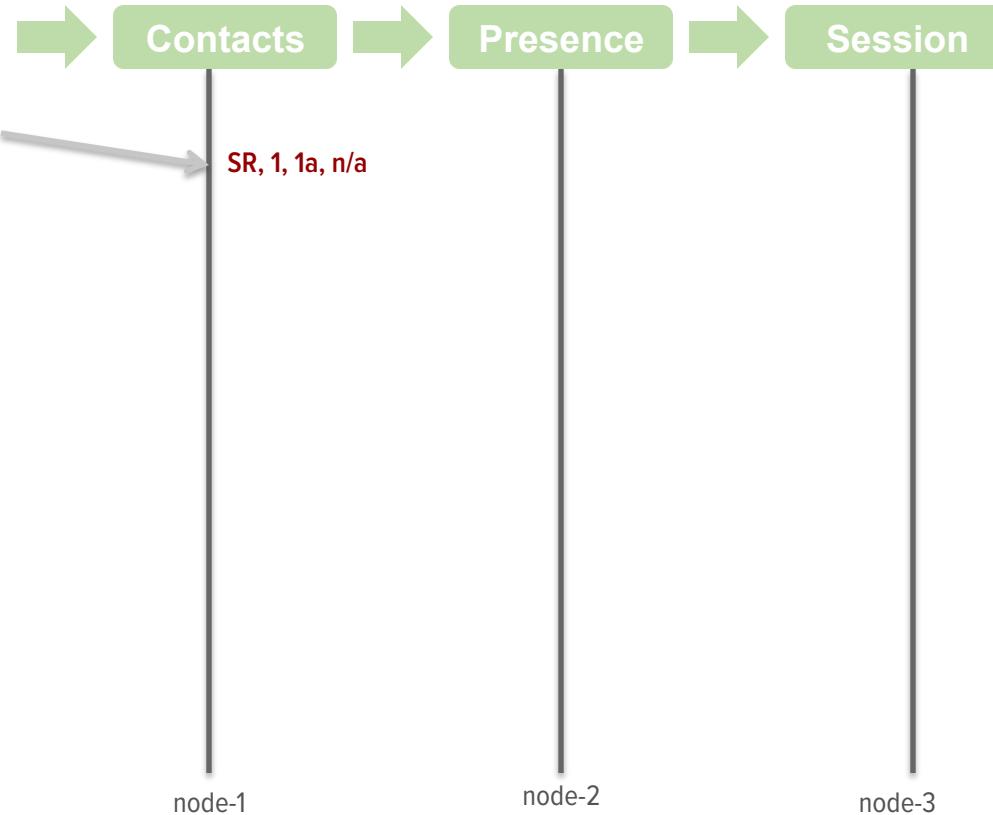
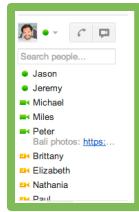
ParentId : 1a

# Review

- Traceld
  - remains unique across systems
- Span
  - can identify RPC attributes
  - example: IP Address between hosts
  - Request/Response

# Warning!

- We are going from 0 to 100 mph very quickly



## Goal

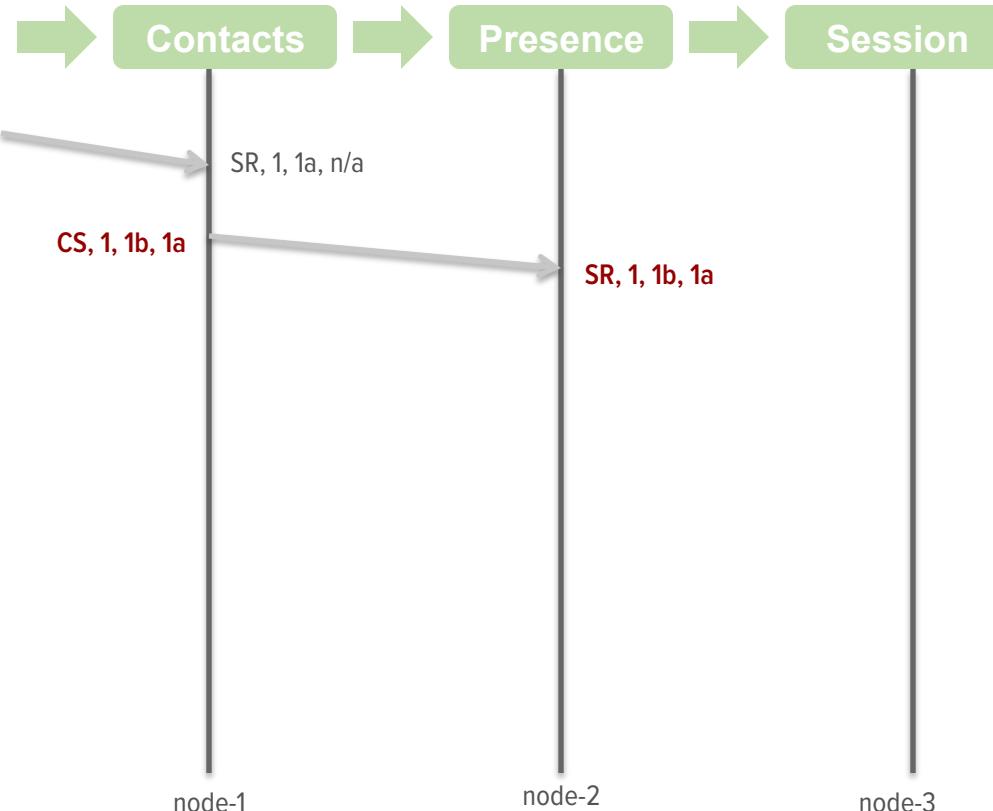
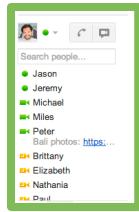
Stitch call tree from local logs

## Node Log View

**node-1:** SR, 1, 1a, n/a

**node-2:**

**node-3:**



## Goal

Stitch call tree from local logs

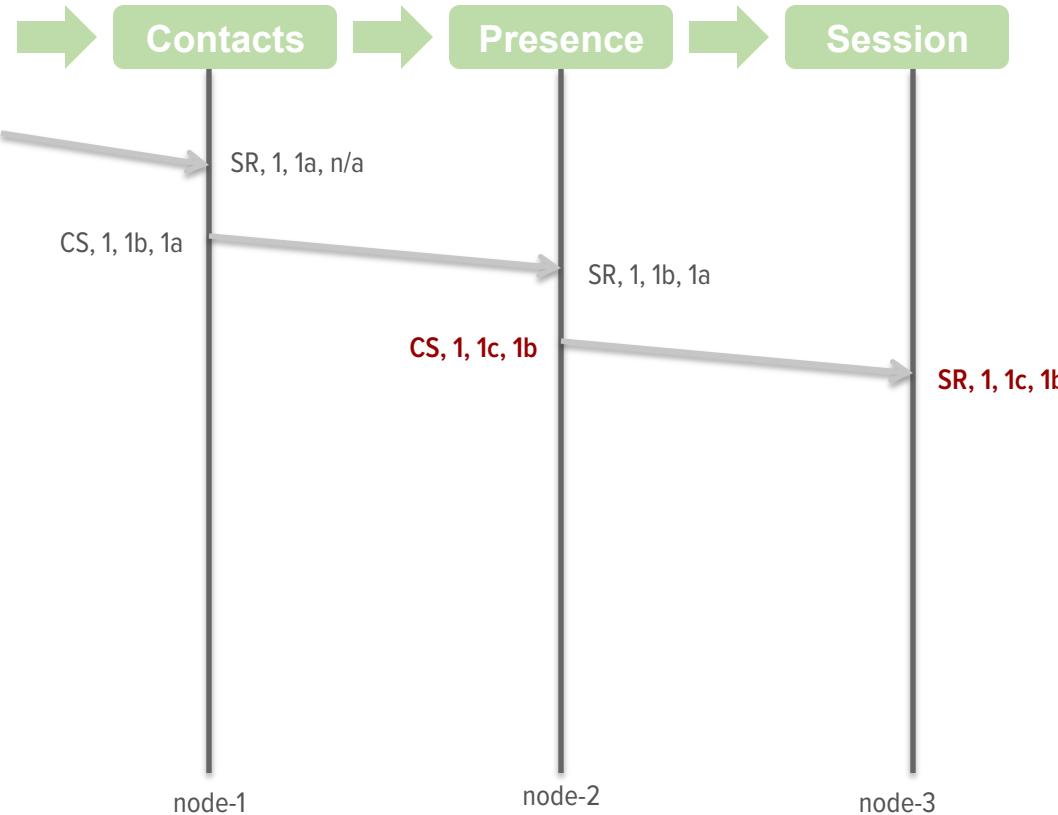
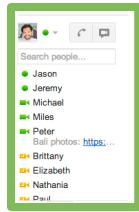
## Node Log View

**node-1:** SR, 1, 1a, n/a

**CS, 1, 1b, 1a**

**node-2:** SR, 1, 1b, 1a

**node-3:**



## Goal

Stitch call tree from local logs

## Node Log View

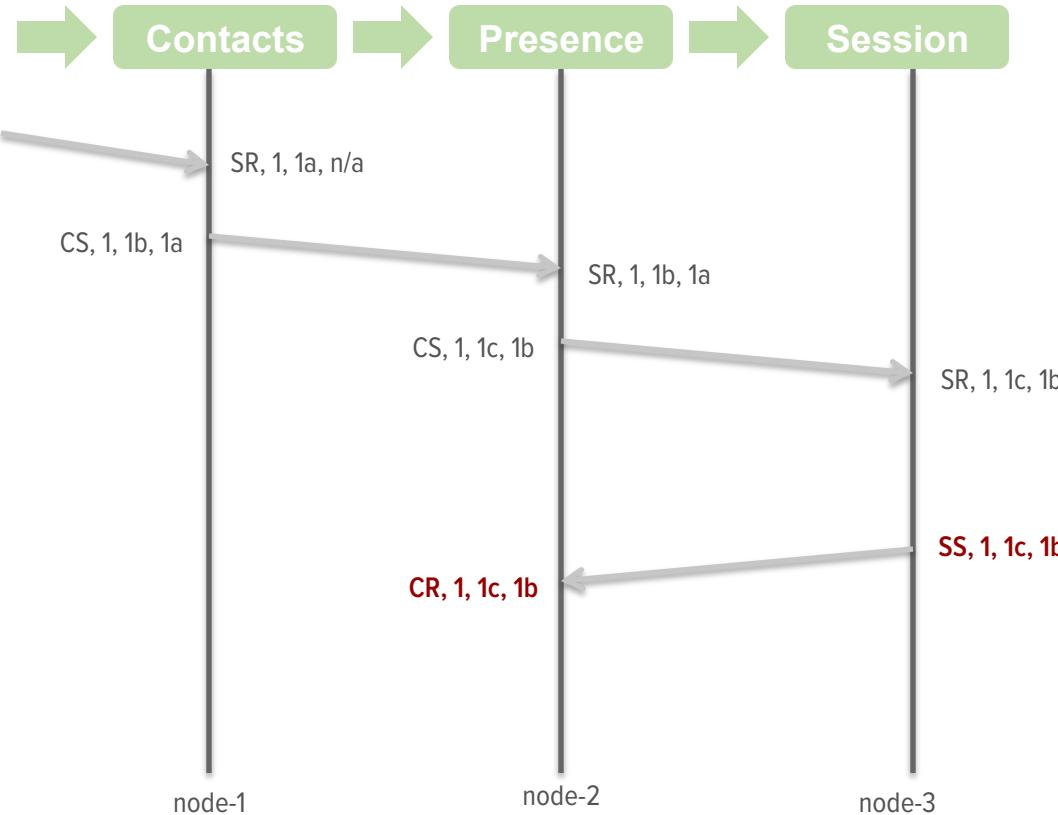
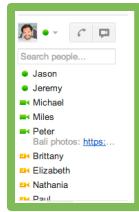
**node-1:** SR, 1, 1a, n/a

CS, 1, 1b, 1a

**node-2:** SR, 1, 1b, 1a

CS, 1, 1c, 1b

**node-3:** SR, 1, 1c, 1b



## Goal

Stitch call tree from local logs

## Node Log View

**node-1:** SR, 1, 1a, n/a

CS, 1, 1b, 1a

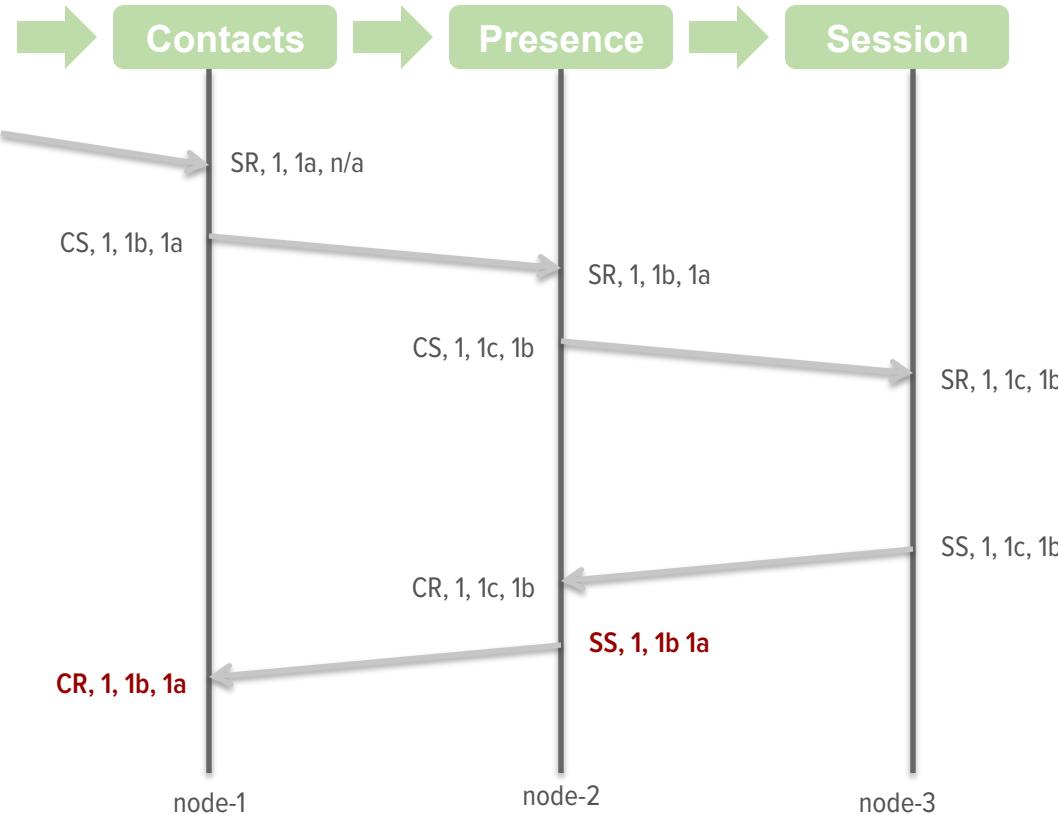
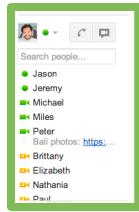
**node-2:** SR, 1, 1b, 1a

CS, 1, 1c, 1b

CR, 1, 1c, 1b

**node-3:** SR, 1, 1c, 1b

SS, 1, 1c, 1b



## Goal

Stitch call tree from local logs

## Node Log View

**node-1:** SR, 1, 1a, n/a

CS, 1, 1b, 1a

**CR, 1, 1b, 1a**

**node-2:** SR, 1, 1b, 1a

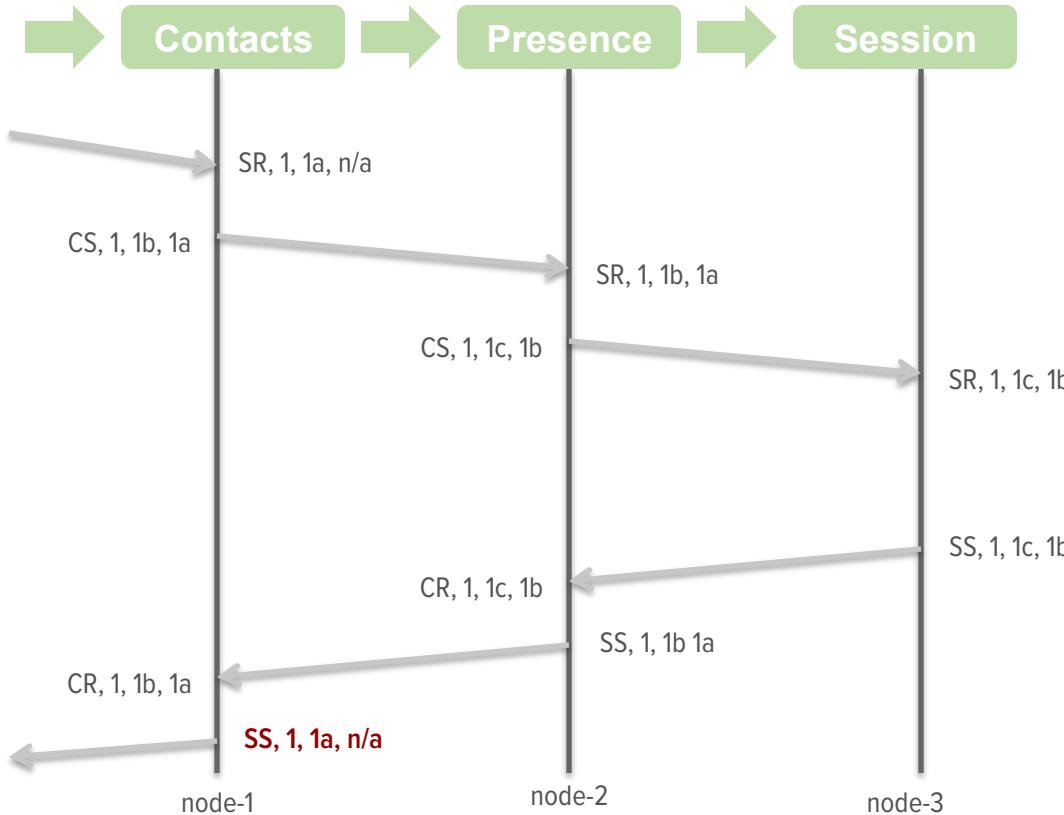
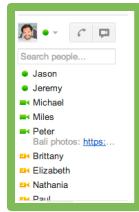
CS, 1, 1c, 1b

CR, 1, 1c, 1b

**SS, 1, 1b, 1a**

**node-3:** SR, 1, 1c, 1b

SS, 1, 1c, 1b



## Goal

Stitch call tree from local logs

## Node Log View

**node-1:** SR, 1, 1a, n/a

CS, 1, 1b, 1a

CR, 1, 1b, 1a

**SS, 1, 1a, n/a**

**node-2:** SR, 1, 1b, 1a

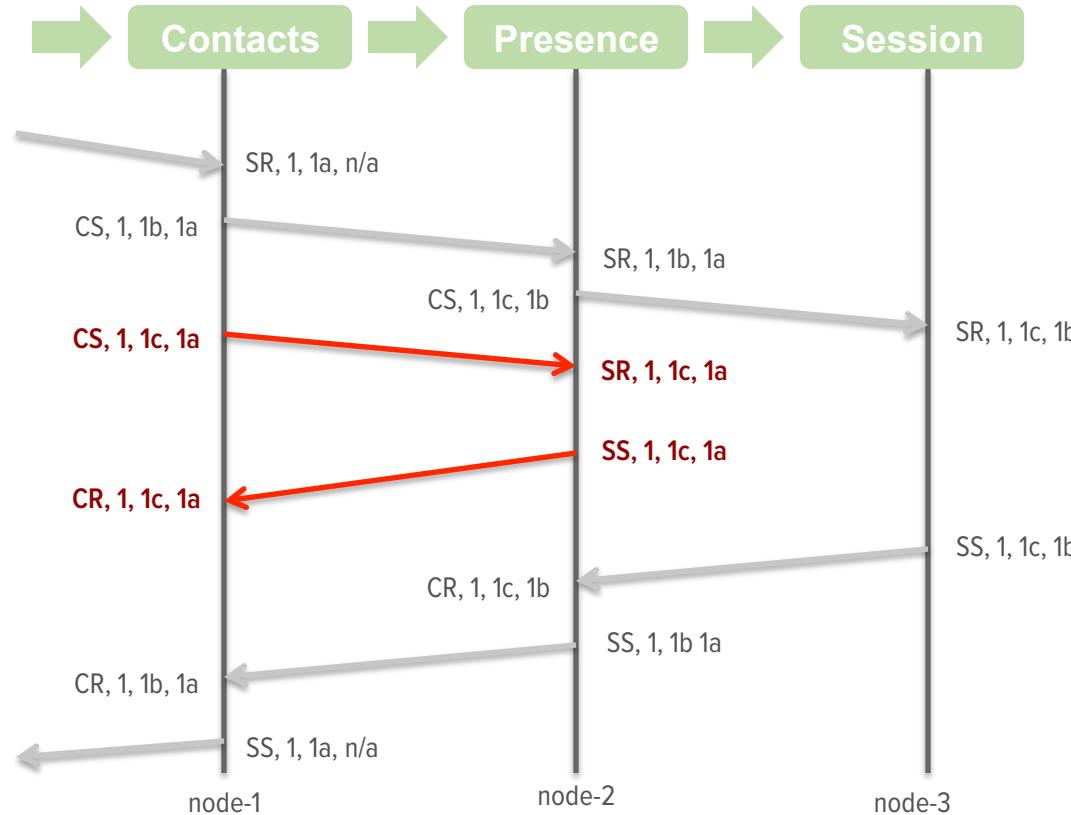
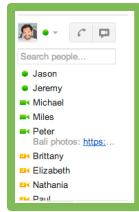
CS, 1, 1c, 1b

CR, 1, 1c, 1b

**SS, 1, 1b, 1a**

**node-3:** SR, 1, 1c, 1b

SS, 1, 1c, 1b



## Goal

Stitch call tree from local logs

## Node Log View

**node-1:** SR, 1, 1a, n/a  
CS, 1, 1b, 1a  
CR, 1, 1b, 1a  
**CS, 1, 1c, 1a**  
**CR, 1, 1c, 1a**  
SS, 1, 1a, n/a

**node-2:** SR, 1, 1b, 1a  
CS, 1, 1c, 1b  
CR, 1, 1c, 1b  
**SR, 1, 1c, 1a**  
**SS, 1, 1c, 1a**  
CR, 1, 1c, 1b  
**CR, 1, 1c, 1a**  
**SS, 1, 1b, 1a**

**node-3:** SR, 1, 1c, 1b  
SS, 1, 1c, 1b

# Review

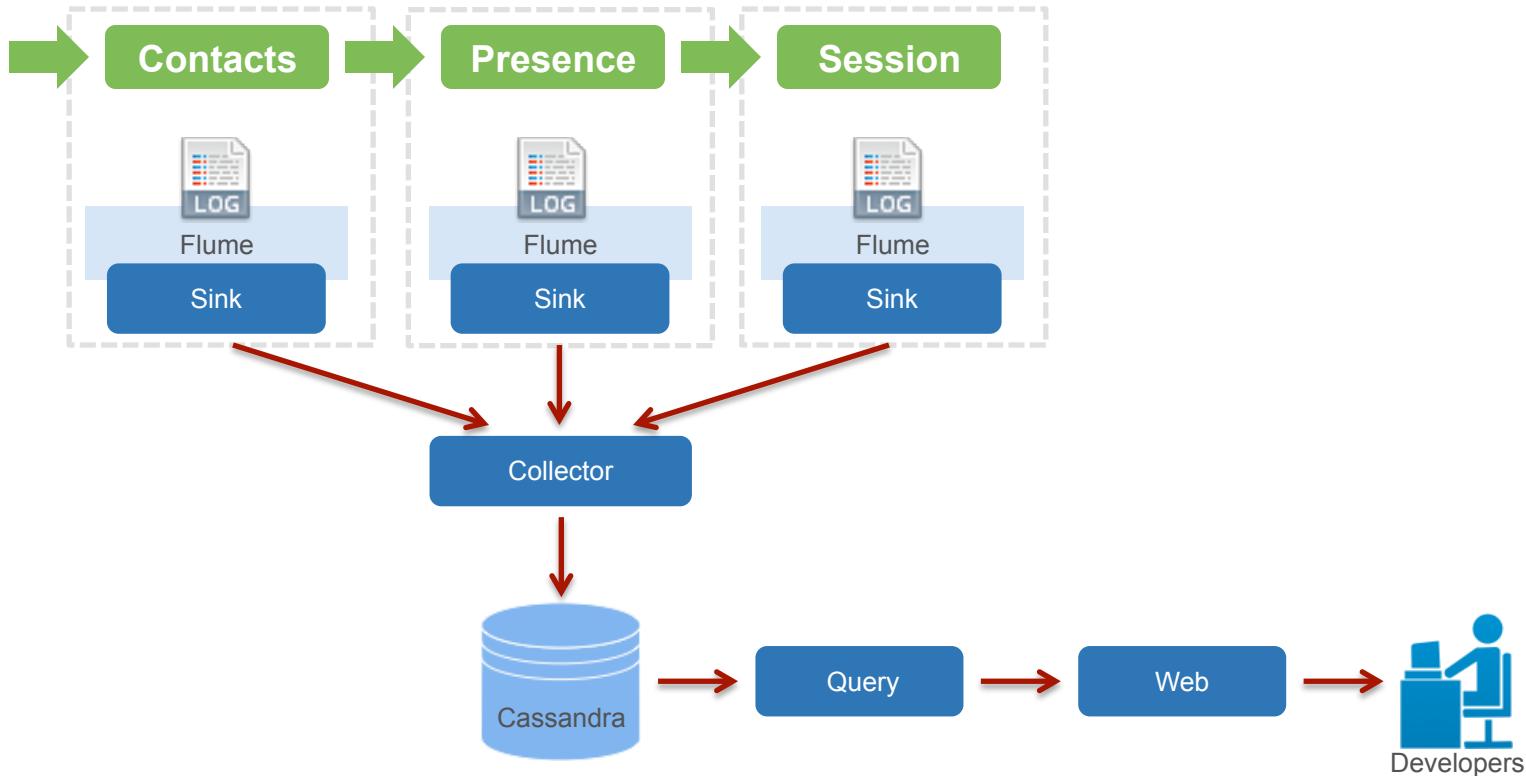
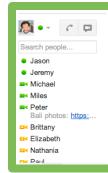
- Annotation based instrumentation
  - Pros
    - End-to-end per request visibility
    - Dependency of a single service on upstream and downstream services
    - Critical path or long tail latency can be identified
  - Cons
    - Requires instrumentation of programs

# Trace collectors

- Out of band trace collection
  - In band trace data adds application overhead
  - Many systems continue munging data after results are returned
  - Kafka, Flume, Scribe collectors are most popular
  - Syslog-ng requires customization

# Trace collection overhead

- PLT test
  - collect empirical data
  - tune sampling
  - second round sampling on collection
- Ability to switch off tracing
  - No impact to in band services



# Demo

- Generate Traces
- Aspects of Zipkin UX
- Aggregate Services Mapped

# Anatomy of a Trace

- Trace
  - root span relating to many spans
- Span
  - set of annotations
- Annotation
  - records time of event + custom
- Sampling
  - boolean

# Sampling

- Used to control the overhead
- Accomplish by using Trace Filters
  - Fixed
  - Adaptive
  - Debug

# Annotation

- Timestamp events
- SR – server received, SS – server sent
- CS – client sent, CR – client received

# Annotation

```
# some event took place, either one by the framework or by the user
struct Annotation {
    1: i64 timestamp          # microseconds from epoch
    2: string value           # what happened at the timestamp?
    3: optional Endpoint host # host this happened on
    4: optional i32 duration   # how long did the operation take? microseconds
}
```

# Annotation.Endpoint

```
# this represents a host and port in a network
struct Endpoint {
    1: i32 ipv4,
    2: i16 port
    3: string service_name # which service did this operation happen on?
}
```

# Span

- Traceld
- name
- Id
- ParentSpanId
- Annotation List
- Binary Annotation List

# Span

```
struct Span {  
    1: i64 trace_id, # unique trace id, use for all spans in trace  
    3: string name, # span name, rpc method for example  
    4: i64 id, # unique span id, only used for this span  
    5: optional i64 parent_id, # parent span id  
    6: list<Annotation> annotations, # list of all annotations/events that occurred  
    8: list<BinaryAnnotation> binary_annotations # any binary annotations  
}
```

# Communicate a Trace

- Protocol headers
  - X-B3-Traceld
  - X-B3-SpanId
  - X-B3-ParentSpanId
  - Sampling
  - ~426 bytes

# Instrumentation Points

- Each framework maybe different
- The ones implemented
- Spring
  - Events, proxies and interceptors
- Netty
  - Message headers, EventBus
- Loggers
  - Logback, Log4j

# Spring Integration

- AbstractReplyProducingMessageHandler
  - service-activator
  - handleRequestMessage
    - inInterceptor
    - handleMessage
    - outInterceptor

# Implementing a Library

- To record SR/SS timestamp annotations
  - `def preProcess(service: String, headers: Map, annot: Map)`
  - `def postProcess()`
- To record CS/CR timestamp annotations
  - `def preClientExec(service: String, headers: Map, annot: Map)`
  - `def postClientExec()`
- Record traces
  - `def trace(message: String)`

# Initialize your System

- List<SpanCollector> spanCollector
  - example: Log4j, Kafka, Scribe, Syslog
- List<TraceFilter> traceFilters
  - example: FixedRate, Adaptive, Debug
- ServerTracer, ClientTracer
  - spanCollector and traceFilter

# Pre-process API

- If part of a Trace
  - set state from current trace headers
  - generate a new span
- If not part of a Trace
  - generate a unique traceld
  - may set spanId = traceld
- set serverReceived (SR)

# Post-process API

- add annotations
  - host, sessionid, jvm id, threadId
- set serverSend (SS) timestamp annotation
- Operations are symmetrical
  - preClientExec
  - postClientExec

# TraceId generation

```
/**  
 * This salt is used to prevent traceId collisions between machines.  
 * By giving each system a random salt, it is less likely that two  
 * processes will sample the same subset of trace ids.  
 */  
private val salt = new Random().nextLong()
```

- Rejoinder
  - ..when you have no colocated services

# TraceId generation

- UUID based
  - <https://github.com/cogitate/twitter-zipkin-uuid>
  - Util
  - Finagle
  - Ostrich
  - Server

# TraceId generation

- UUID support
  - SQL Lite
  - Cassandra
- Circa April 2015

# Useful Zipkin APIs

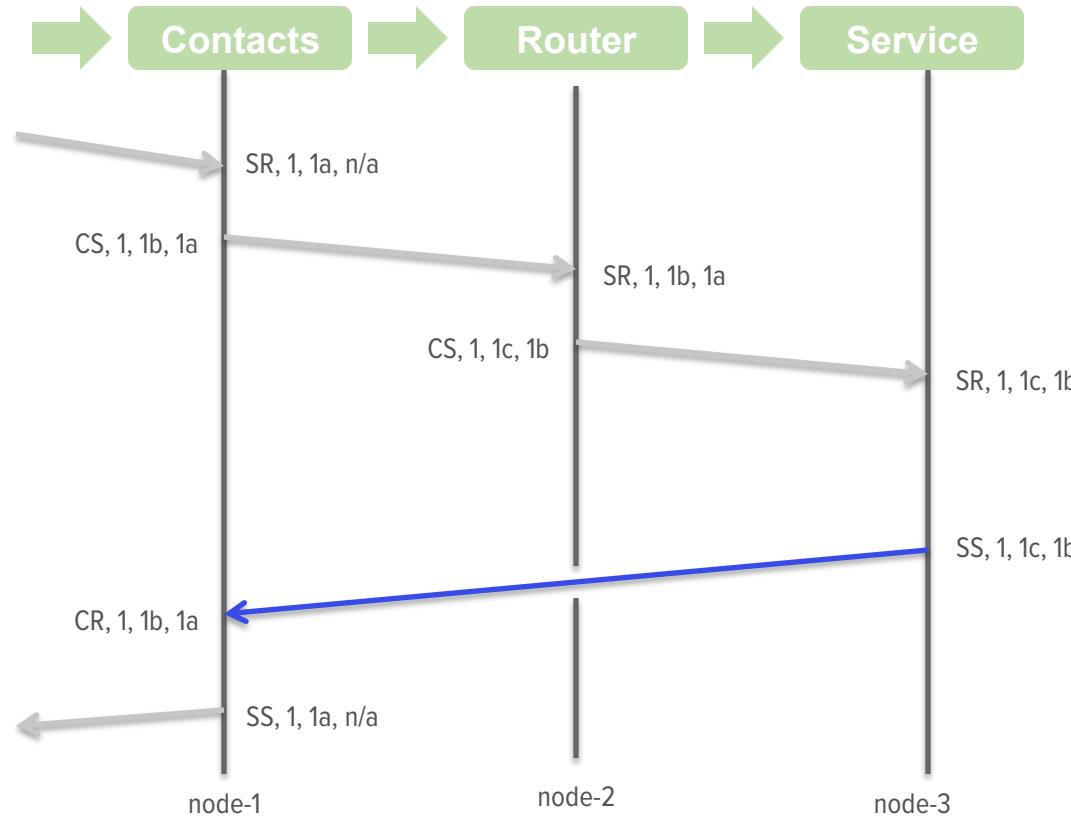
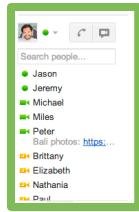
- /traces/{id}
- /api/trace/{id}
- /api/query
- /api/services
- /api/spans

# Production coverage

- Common RPC libraries provide leverage
- Some coverage maybe lost due to
  - non-instrumentable legacy systems
  - non-standard control flow

# Non standard use cases

- Routers and Forwarders
- Netty like dispatchers
- Non-participating trace nodes
  - To Trace or Not to Trace



## Goal

Stitch call tree from local logs

## Node Log View

**node-1:** SR, 1, 1a, n/a

CS, 1, 1b, 1a

CR, 1, 1b, 1a

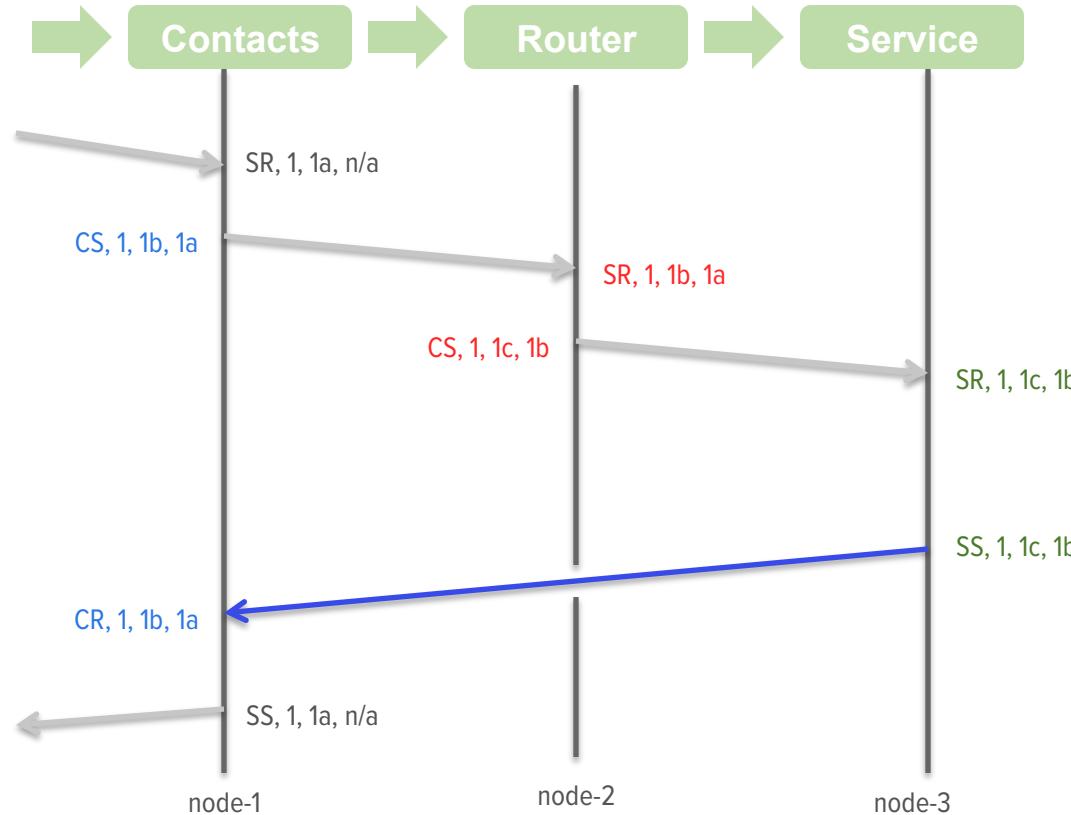
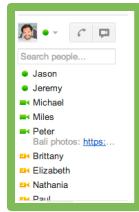
SS, 1, 1a, n/a

**node-2:** SR, 1, 1b, 1a

CS, 1, 1c, 1b

**node-3:** SR, 1, 1c, 1b

SS, 1, 1c, 1b



## Goal

Stitch call tree from local logs

## Node Log View

**node-1:** SR, 1, 1a, n/a

CS, 1, 1b, 1a

CR, 1, 1b, 1a

SS, 1, 1a, n/a

**node-2:** SR, 1, 1b, 1a

CS, 1, 1c, 1b

**node-3:** SR, 1, 1c, 1b

SS, 1, 1c, 1b

Where's my span?

# Non standard use cases

- Reactive, Netty, Asynchronous callbacks
  - Request sent on a different thread
  - Response received on a different thread
- Messaging pattern
  - `getSpan` (get context)
  - `setSpan` (set context)

# Lessons Learnt

- Development teams use tracing system to debug
  - Surprised with sampling
- Requirement to generate TracId ALWAYS
- Multipage group tracing
  - UberTrace – List<TracId>
- Use collectors that you already have
  - add converters
- Monitor
  - Disk Space
  - Collectors buffering

# Lessons Learnt

- Dependency Hell
  - control libthrift version
- Mis-understand cassandra
  - CQL
  - Cassandra-CLI
- Need for different consumers
  - Elastic Search
  - Hadoop
- Aggregations
  - Scalding, Algebird, Spark
  - Steep learning curve

# Lessons Learnt

- Be ready to interact with INFRA
  - Switch/Port, Dynamic Port Info
  - SSD vs Traditional Hard disk
  - Tools to probe system information
- Network failure points
- Capacity planning and sizing
- Understand Storage systems
  - constraints and limitations

# Call Tracing

- Dedicated to a full view of a request's path through SOA
- Based on Google's Dapper paper
- Zipkin, Brave, Node-Tryfer
  - Scala, Java, NodeJS
- Zipkin
  - Collector, Query and UI

# Acknowledgements

- SpringOne2GX Team!
- Jeff Smick – Twitter
- Kristof Adriaenssens – Brave
- Node Tryfer - Rackspace
- Eirik Sletteberg - Aggregates

# Stay Connected

<https://www.linkedin.com/in/cogitate>

# Q&A